

Sub-Word Transformer Distillation for Multilingual Low-Resource NLP

Authors: Sunita Verma, R. Deshmukh

Publication: ACM Transactions on Asian and Low-Resource Language Information Processing, 2026

ABSTRACT

We demonstrate a novel knowledge distillation pipeline achieving 94% BERT performance with 65% fewer parameters across 14 Indic languages.

1. INTRODUCTION

Modern cloud-native systems and transformer architectures require stringent verification, fault isolation, and reproducible evaluation benchmarks across national computational grids.

2. SYSTEM ARCHITECTURE & METHODOLOGY

Our empirical evaluation framework deploys automated telemetry probes, tracking p99 request latencies, gradient throughput, and memory pressure during high concurrency stress testing.

3. EXPERIMENTAL RESULTS

Across 10,000 synthetic request iterations, our method demonstrated superior stability, achieving an overall 43 percent decrease in tail latency and 94 percent model fidelity on benchmark datasets.

4. CONCLUSION & FUTURE WORK

We established an open-source baseline for researchers across national academic institutions. Future extensions will incorporate federated edge telemetry and multi-modal distillation.